EXHAUSTSHIELD COMMAND CENTER · EXECUTIVE BRIEF

Where the Market Stands vs. Where We Stand

A honest read on our AI data-leak-prevention pilot against the latest global developments in shadow AI governance — and the fastest path to monetizing it.

Prepared Aug 10, 2026 · Market data through Aug 6, 2026

Market Timeline — 12 Months of Shadow AI Governance

2025
Sensitive share of data pasted into AI tools more than triples: 10.7% → 34.8% (Cyberhaven).

SEP 2025
SentinelOne acquires Prompt Security; Cato acquires Aim Security — consolidation begins.

OCT 2025
Check Point acquires Lakera — a second independent folded into a security major.

NOV 2025
Gartner: up to 40% of enterprises will face a shadow-AI-related breach by 2030.

FEB 2026
EU AI Act enforcement begins — inadequate AI data controls now carry real penalties.

JUN 2026
Zscaler: 18,033 TB moved to AI apps in 2025 (+93% YoY); 410M ChatGPT DLP violations.

AUG 4, 2026
Sentra launches Shadow AI DLP — a funded DSPM leader enters our exact category. See p.2.

Market Today vs. Our Pilot

MARKET TODAY — What everyone else ships

- ✕ **Visibility and gating.** Tools report what left, or decide whether it goes — block, allow, log, warn, or ask for justification.
- ✕ **Enforcement stops at the decision.** Once a prompt is flagged, the common pattern is block or pass — not editing the content itself.
- ✕ **Fragmented by surface.** Browser extension here, endpoint agent there — few unify gateway + browser + agent traffic.
- ✕ **Vendor churn.** Six major acquisitions in 12 months (Lakera, Prompt Security, Aim Security, SurePath AI, LayerX, Wiz) — stability itself is a live buyer risk.

OUR PILOT — What we already built

- ✓ **Three-lane gate.** Safe / Medium / High classification, with Medium prompts hidden or rewritten before they leave.
- ✓ **Fail-closed by design.** Hard secrets never leave; uncertainty always escalates up, never quietly down.
- ✓ **Local knowledge answers.** Policy Q&A comes from a vetted vault — zero AI call, zero leak risk.
- ✓ **Human correction loop.** Disputes and corrections are captured — and never used to train a model.

99%
of organizations already have sensitive data exposed to AI tools (Varonis, 2026)

77% / 82%
of employees paste data into GenAI — 82% of that via personal, unmanaged accounts (LayerX/IBM)

$670K
added to average breach cost when shadow AI is involved (IBM, 2025)

$2.1B → $15.5B
AI data-readiness / shadow-AI-DLP market, 2026 → 2030 (Sentra)

THE NEWEST DEVELOPMENT

A Funded DSPM Leader Just Entered This Exact Space

Four days before this brief, Sentra launched a direct competitor to the category we're building in. Here's what it means for us.

AUG 4, 2026 · GENERAL AVAILABILITY

Sentra — Shadow AI DLP

- Detects AI-app use and stops sensitive data from reaching ChatGPT, Claude, Gemini, or Copilot — a new module inside Sentra's existing Data Readiness Platform.
- Sentra sizes the category at $2.1B in 2026, growing to $15.5B by 2030.
- Runs on Sentra's existing petabyte-scale data-security customer base — real data, not a fresh start.
- Cites the same pain point we target: data, not models, is the biggest obstacle to scaling AI.
- Its published controls are discovery, plus block, allow, log, justify, or verify — five gate-style actions on whether a prompt goes through, not on what's inside it.

What This Means for Us — Read It Straight

- ✓ Validation. A funded DSPM vendor entering this space confirms the problem is real and buyers are moving now.
- ✓ Their public feature list is gate-only. Nothing in Sentra's launch materials describes editing or rewriting a prompt's content — only deciding whether it goes.
- ✗ They have what we don't: petabyte-scale data classification and an existing enterprise customer base.
- ✗ Competing head-to-head as a discovery platform isn't realistic at our stage.
- ✓ Our lane: the layer that keeps work moving, sitting in front of or alongside discovery tools like Sentra's.

Strategic read: This is a read of Sentra's public launch materials, not a hands-on product test — feature lists can undersell what's shipped. Treat it as "not documented," not "not possible." Sentra competes on data-classification infrastructure; our clearest opening is the moment after detection — what happens to the request next.

CONFIRMED ADVANTAGES

Two Things We Checked For — and Didn't Find Elsewhere

Based on public product materials for Nightfall, Harmonic Security, WitnessAI, and Sentra as of Aug 2026. Absence from public materials is evidence, not proof — vendors don't always publish everything they ship.

1. Content rewriting, not just gating

- ✓ Our Medium lane rewrites a risky prompt into a safer version of the same question, so it can still be sent and answered.
- ✓ Across Sentra, Nightfall, Harmonic, and WitnessAI's published controls — block, allow, log, warn, coach, justify, verify — none describe editing prompt content.
- ✕ Confidence: moderate. Based on public feature lists, not a lab test of each product.

2. A local answer layer with human review

- ✓ Company-policy questions are answered from a vetted local vault — no AI call at all — with a dispute/correction workflow that updates the vault but never trains a model.
- ✓ None of the four vendors reviewed publish an equivalent. WitnessAI comes closest by routing sensitive queries to an internal model — still an AI call, not a vault lookup.
- ✕ Confidence: higher. This isn't a missing feature description — it's a different category of feature none of them market at all.

How We Tap This

LEAD WITH #2

The knowledge vault has no visible competitor. It's the safer claim to put in front of a buyer or a partner conversation.

POSITION #1 CAREFULLY

Frame rewriting as "the step after detection that block-only tools skip" — not as a claim no competitor could ever build.

SELL THE PAIR

Together they're a single pitch: don't just stop the prompt, answer it safely — from a vault, or with the sensitive parts removed.

CAPABILITY MATRIX & THE PATH TO PRODUCT

Where We Already Win, Where It's Still a Pilot

Capability	Market Tools	Sentra (newest)	Our Pilot	Why It Matters
Risky prompt	Block / coach / warn	Block, allow, log, justify, verify	Hides / rewrites, task continues	Confirmed advantage — see p.3
Company Q&A	Out of scope	Out of scope	Local vault, no AI call	Confirmed advantage — see p.3
Hard secrets	Varies	Blocks sharing	Fail-closed, private AI only	Contained, not just refused
Data classification	Varies	Petabyte-scale	Rule-based, single gateway	Our clearest gap
Audit trail	Inconsistent	Platform-level	Local ledger (pilot)	Needed for trust
Enterprise auth	Standard	Enterprise-grade	Not yet built	Required to sell
Agent traffic (MCP)	Emerging	Not stated	Roadmap	Next battleground

Two Ways to Use It Today

INTERNAL

Run the pilot on our own AI traffic now — the gate logic already works.

EXTERNAL · PHASE 2

Sell as the continuation and answer layer, not a discovery tool — pairs with Sentra, Nightfall, Harmonic.

FASTEST WIN

Point the gateway at real traffic. Near-zero cost, immediate value.

Honest Gaps Before It's a Product

COVERAGE

Add browser + endpoint capture — the edge competitors already have.

TRUST & SECURITY

Tamper-evident audit, SSO, RBAC, multi-tenancy — required before selling.

DETECTION DEPTH

Add ML classification to cut false positives at scale.

Bottom line: Sentra just entered our space with data-classification muscle. Its published controls don't include rewriting content or answering from a local vault — two things our pilot already does, based on public materials as of Aug 2026. What's missing on our side isn't the idea, it's infrastructure: coverage, auth, and a tamper-evident audit trail. Deploy internally now; position externally as the layer that plugs alongside tools like Sentra's, not against them.