Your prompts are now a priced asset

Three options to govern employee AI use, costed end to end for 500 – 3,000 people. Every figure is a published list price or a named published study.

The trigger: Meta now publishes what your data is worth

Meta's Muse Spark 1.3 rate card: **$1.25 / $4.25** per million tokens on Standard (never train) vs **$0.10 / $0.20** on Contributor — a **92–95% discount** for the right to train on your prompts. That spread is about **$1.24 per million tokens** of business prompt data. Our employees are giving that away for free today.

What the market pays for content and data

Buyer → seller	Published value	Why it matters
Meta → News Corp	Up to **$50M / year** (3-yr, ~$150M)	Proprietary text has a market price
OpenAI → News Corp	**$250M** over 5 years	Largest known single-publisher deal
Google → Reddit	**~$60M / year**	Even user chatter is monetised
Meta → Scale AI	**$14.3B** for 49% stake (2025)	Data pipelines valued in billions
Meta → you	**$1.24 / 1M tokens** discount	**They will pay to read our prompts**

Meanwhile, this is what leaves the building today

11%
of pastes into ChatGPT are sensitive · **Cyberhaven**

39.7%
of AI interactions involve sensitive data · **Cyberhaven**

82%
of GenAI pastes from unmanaged accounts · **LayerX**

$19.5M
avg annual insider-risk cost · **Insider Risks 2026**

Finding	Number	Source
GenAI is #1 route for company→personal data movement	**32%** of that exfiltration	LayerX
Orgs that found staff using unapproved AI with company data	**65%** in last 12 months	Kiteworks 2026
Orgs with AI-specific DLP actually deployed	**under 30%**	Kiteworks 2026
Outright AI bans collapsed — without controls replacing them	**28% → 7%**	Cisco 2026
How often an employee feeds sensitive data to AI	**once every 3 days**	Cyberhaven 2026

Board question: when up to 3,000 people paste company text into models that bid for training rights, who decides what is allowed to leave? Three options follow.

How every number is calculated

One formula · three usage scenarios · published rate cards. No blended estimates.

Published rate cards (verified Sep 2026)

Line item	List price	Line item	List price
M365 E3	**$39**/user/mo (from Jul 2026)	**M365 E5**	**$60**/user/mo (from Jul 2026)
Copilot add-on	**$30**/user/mo annual	**Purview Suite (E3)**	~**$12**/user/mo
Studio PAYG	**$0.01**/credit	**Studio pack**	**$200** / 25k credits/mo
Groq gpt-oss-120b	**$0.15** in / **$0.60** out /1M	**Azure GPT-4o-mini**	**$0.15** in / **$0.60** out /1M
Azure GPT-5-mini	**$0.50** in / **$2.00** out /1M	**Azure GPT-5**	**$5.00** in / **$20.00** out /1M
Azure VM D2s_v5	**$70.08**/mo East US Linux	**Studio credits**	classic 1 · gen 2 · action 5 · graph 10

Cost of one gated chat = 1,000 input + 500 output tokens

Model	Input	Output	Per chat
Groq gpt-oss-120b	0.001M × $0.15 = $0.00015	0.0005M × $0.60 = $0.00030	**$0.00045**
Azure GPT-4o-mini	same as Groq	same as Groq	**$0.00045**
Azure GPT-5-mini	0.001M × $0.50 = $0.00050	0.0005M × $2.00 = $0.00100	**$0.00150**
Azure GPT-5 flagship	0.001M × $5.00 = $0.00500	0.0005M × $20 = $0.01000	**$0.01500**

Groq and Azure GPT-4o-mini are priced identically — model choice is quality, not lock-in.

Three usage scenarios (22 working days / month)

Scenario	Chats / user / day	500 users	1,500 users	3,000 users
BEST — light pilot	1	11,000 / mo	33,000 / mo	66,000 / mo
NORMAL — steady	3	33,000 / mo	99,000 / mo	198,000 / mo
WORST — heavy	8	88,000 / mo	264,000 / mo	528,000 / mo

Conservative: tables below bill 100% of chats to a paid model. ~4 in 10 are local policy answers with **no model call**, so live bills come in lower.

OPTION 1 — Stay with existing Copilot (few E5 · majority E3 + basic Copilot Chat)

No build. No new dashboard. Governance = whatever that licence tier already gives.

Employee — Copilot Chat — or quietly ChatGPT / Claude in a browser

→ **Copilot Chat** — Included with E3/E5 · Microsoft-hosted

→ **Purview DLP** — SIT match → block prompt or block web grounding

COVERED	NOT COVERED	NOT COVERED	NOT COVERED
Prompt inside Copilot	Same person on chatgpt.com	No scrub-then-send	No private model for secrets

Option 1 — cost and verdict

Cost line	500 users	1,500 users	3,000 users
Licences already paid **20% E5 @$60 + 80% E3 @$39**	$21,600 / mo **$259,200 / yr**	$64,800 / mo **$777,600 / yr**	$129,600 / mo **$1,555,200 / yr**
New spend to keep as-is	**$0**	**$0**	**$0**
Purview depth for E3 majority **~$12 × E3 headcount**	**$57,600 / yr**	**$172,800 / yr**	**$345,600 / yr**
Full Copilot for everyone **$30 × headcount**	**$180,000 / yr**	**$540,000 / yr**	**$1,080,000 / yr**

Strength	Limitation
Zero new spend, zero build, familiar in Office	Controls apply **only inside Copilot**. ChatGPT = no gate, no log. Purview blocks — cannot scrub-then-answer. No private-model lane.

OPTION 2 — ExhaustShield gate + Microsoft brain (Azure OpenAI instead of Groq)

Our four lanes decide what may leave · approved text goes to a Microsoft-hosted GPT model

Any employee
Any licence · one URL or Teams tab

→

ExhaustShield gate
Scan → classify → confirm / scrub / hold

→

Azure OpenAI
Our tenant · Key Vault · no Copilot seat

SAFE
Clean → send after confirm

MEDIUM
Hide values → scrubbed send

HIGH
Keys / PHI → private model only

LOCAL
Policy → our files, no model

How to implement in the enterprise

#	Action
1	**Host** — deploy gateway to internal VM / App Service with TLS
2	**Identity** — Entra SSO so every prompt is attributable (E3, E5, or no Copilot)
3	**Model** — Azure OpenAI resource · key in Key Vault · point gate external-AI setting at it
4	**Lanes** — Safe + scrubbed Medium → Azure OpenAI · High stays private · Local never calls a model
5	**Assurance** — InfoSec test: High-lane text never appears in any Azure OpenAI request log

Cost — NORMAL scenario (3 chats / user / day), all chats billed

Model behind gate	500 · 33k chats	1,500 · 99k chats	3,000 · 198k chats
Azure GPT-4o-mini @ $0.00045	$14.85 / mo	$44.55 / mo	$89.10 / mo
Azure GPT-5-mini @ $0.00150	$49.50 / mo	$148.50 / mo	$297.00 / mo
Hosting (D2s_v5 × instances)	$70 / mo	$140 / mo	$280 / mo
Total with GPT-5-mini	**$120 / mo · $1,434 / yr**	**$289 / mo · $3,462 / yr**	**$577 / mo · $6,924 / yr**

Caution: routing through the real Copilot Chat brain costs ~100× more

Copilot Chat API / Work IQ needs a **per-user M365 Copilot licence** — **$180k / $540k / $1.08M** a year at 500 / 1,500 / 3,000 people. **Azure OpenAI = GPT-class quality, no per-seat fee.**

OPTION 3 — Standalone ExhaustShield dashboard (our pilot, as built)

One front door for everyone · five tabs · Groq as low-cost external brain · runs today

Any employee
Intranet link or Teams website tab · Entra SSO

→

5-tab dashboard
Overview · Knowledge · Gate · Vault · Lab

→

Four lanes
Safe / Medium / High / Local · every ask logged

→

External brain
Groq (or swap to Azure OpenAI later)

SAFE
Groq after user confirms

MEDIUM
Scrubbed text only → Groq

HIGH
Private Ollama · never external

LOCAL
Local files or human vault

How to implement in the enterprise

#	Action	Owner
1	**Host** — existing gateway onto one internal VM with HTTPS	IT infra
2	**Identity** — Entra SSO; all 3,000 can sign in, no Copilot seat needed	Identity
3	**Brain** — Groq key in server env (or Azure OpenAI later — same setting)	Platform
4	**Private lane** — Ollama on same or second VM for High lane	Platform
5	**Reach** — publish as intranet link + Teams website tab	Comms / IT
6	**Evidence** — Overview tab: sensitive words, lanes, team breakdown	Security

Cost across BEST / NORMAL / WORST (Groq @ $0.00045 + hosting)

Scenario	500 users	1,500 users	3,000 users
BEST — 1 chat/day **Groq + hosting**	$4.95 + $70 **$900 / yr**	$14.85 + $140 **$1,858 / yr**	$29.70 + $280 **$3,716 / yr**
NORMAL — 3 chats/day **Groq + hosting**	$14.85 + $70 **$1,018 / yr**	$44.55 + $140 **$2,215 / yr**	$89.10 + $280 **$4,429 / yr**
WORST — 8 chats/day **Groq + hosting**	$39.60 + $70 **$1,315 / yr**	$118.80 + $140 **$3,106 / yr**	$237.60 + $280 **$6,211 / yr**

Worst case at 3,000 people ≈ $6,200 / year

528,000 governed conversations a month — roughly **two Copilot seats**. Vs Option 1: covers all staff + four lanes. Vs Option 2: same gate, lower bill today, swap to Azure OpenAI later.

Final comparison — net new annual spend (NORMAL · 3 chats / user / day)

	OPTION 1 Existing Copilot	OPTION 2 Gate + Azure OpenAI	OPTION 3 Standalone dashboard
Net new — 500	$0 · +$57.6k Purview depth	**$1,434**	**$1,018**
Net new — 1,500	$0 · +$172.8k Purview depth	**$3,462**	**$2,215**
Net new — 3,000	$0 · +$345.6k Purview depth	**$6,924**	**$4,429**
Worst case — 3,000	Unchanged ($0 new)	~$10,900	**~$6,211**
Covers non-Copilot staff	No	**Yes**	**Yes**
Blocks / scrub / private / local	SIT block only	**All four**	**All four**

Recommendation · board ask · sources

	OPTION 1	OPTION 2	OPTION 3
Sensitive words + lanes by team	Adoption reports only	**Yes**	**Yes**
Answer quality	GPT-class	**GPT-class**	Open models (swappable)
Training-data control	Vendor terms	**We decide per lane**	**We decide per lane**
Time to live	Live now	3 – 5 weeks	**1 – 2 weeks (built)**
Board risk if we do nothing	**High** — 82% unmanaged	Low	Low

RECOMMENDATION — Adopt Option 3 now · keep Option 1 · hold Option 2 as upgrade path

Dashboard for every employee: **~$1,000 – $4,400 / year** (up to ~$6,200 worst at 3,000). Keep Copilot + Purview for Office — complementary, not a substitute. Want GPT-class behind the gate? Switch to **Azure OpenAI** (~+$2,500 / year) — never buy Copilot seats for everyone at **$180k – $1.08M**.

Three sentences for the minutes

#	Minute
1	Meta discounts tokens **92–95%** for training rights (~**$1.24 / 1M tokens**). Today we hand that over for nothing.
2	**82%** of AI pastes are unmanaged; **11%** of paste content is sensitive — Copilot-only covers a minority of traffic.
3	A gate for all **3,000** people ≈ **$4,400 / year** (NORMAL) — less than **13 Copilot seats**.

Decision requested

Decision	Detail
Approve Option 3 pilot	Deploy org-wide or phased 500 → 1,500 → 3,000. Budget **~$1k – $6.2k / year**.
Keep Option 1 as-is	Keep Copilot / Purview for Office. Do not rely on them alone for ChatGPT / Claude traffic.
Park Option 2 upgrade	If quality complaints arise, switch gate to Azure OpenAI (~+$2.5k / yr). Do **not** buy Copilot for all 3,000.
Success (90 days)	% of AI chats through the gate · High holds · scrubbed Medium sends · zero High text in external logs.

Sources

Microsoft 365 (E3 $39, E5 $60 from Jul 2026; Copilot $30). Copilot Studio Licensing Guide Jun 2026 ($200/25k credits; PAYG $0.01). Microsoft Learn credit rates + Purview DLP for Copilot. Purview Suite ~$12/user/mo (E3). Azure OpenAI (GPT-4o-mini $0.15/$0.60; GPT-5-mini $0.50/$2; GPT-5 $5/$20 per 1M). Groq gpt-oss-120b $0.15/$0.60. Azure VM D2s_v5 East US $70.08/mo. Meta Muse Spark 1.3 Standard $1.25/$4.25 vs Contributor $0.10/$0.20. Meta–News Corp up to $50M/yr; OpenAI–News Corp $250M/5yr; Google–Reddit ~$60M/yr; Meta–Scale AI $14.3B (49%). Cyberhaven 2026 (11%; 39.7%; every 3 days). LayerX (82% unmanaged; GenAI 32% of corp→personal). Kiteworks 2026 (65%; AI DLP <30%; $19.5M). Cisco 2026 (bans 28%→7%). List prices; EA discounts may vary.