PERSONAL EXECUTIVE REFERENCE

ExhaustShield
CEO Demo Playbook

Precise · Color-coded · Q&A; ready
Plain English · Pros / Cons / Production fixes · File ownership

Goal: Whatever question they throw — answer with clarity and confidence.

Start demo: **start_demo.bat** → open **http://127.0.0.1:8000/**
Reset Demo before the room walks in · Ctrl+F5 after UI updates

CONTENTS

1. One-minute story
2. Word list (simple English)
3. Start the demo — checklist
4. 8-minute talk track
5. Pilot vs production
6. Three paths (Safe · Medium · High)
7. Every tab — what to show
8. Scoreboard — how to read numbers
9. Hide & rewrite — without changing the ask
10. Private AI honesty
11. Pros · Cons · Fix in production
12. Killer Q&A; (say these out loud)
13. Demo recovery lines
14. File & folder guide + who speaks
15. What we ship today · What’s next
16. One-page cheat sheet

1 · ONE-MINUTE STORY

ExhaustShield is the security guard in front of AI. It checks every prompt before data can leave. Then it picks a path:

- **Safe** — clean text → confirm → outside AI (e.g. Groq).
- **Medium** — hide/rewrite business secrets → then outside AI.
- **High** — hard secrets → private AI only. Never outside.

Company knowledge answers come from files or humans — not from a guessing model.

SAY: “People will use AI. We stop secrets leaving quietly — and when we don’t know, a human answers.”

2 · WORD LIST — SIMPLE ENGLISH

AI / model / LLM
Chat assistant software that answers in normal language.

Prompt
What someone types into the AI.

Outside AI / Groq
AI on someone else’s computers (internet). Pilot live option = Groq.

Private AI / Ollama
AI on this laptop (or later: company network). High lane uses this.

Lane / path
Risk bucket: Safe, Medium, or High.

Hide / redact
Cover secrets so clear text does not leave.

Rewrite
Same question, safer words (fake names, softer $/%).

Confirm
Pause: “Send outside?” before anything leaves.

Last-mile re-scan
Final auto-check right before send.

Fail-closed
If unsure or private AI is down → block outside send. No leak.

Simulate mode
Fake outside reply for demo; gate decisions still real.

Human vault
Approved answers written by people. Not used to train AI.

RAG
Find docs, then ask AI. Normal RAG often sends clear text; we add a gate.

Egress
Data leaving to the public internet / vendor.

Evidence / audit
Saved proof of what happened and why.

Reset Demo
Clears scoreboard + evidence. Keeps human vault on purpose.

127.0.0.1 / localhost
This computer — not a public website.

3 · START THE DEMO

- Optional: start Ollama + model **llama3.2** (needed for High answers).
- Optional: real key in **groq_key.local.bat** for live outside replies.
- Run **start_demo.bat** → open **http://127.0.0.1:8000/**.
- Ctrl+F5 if UI looks stale · click **Reset Demo** for a clean board.

SAY UP FRONT: "Outside AI is live" or "Outside AI is simulated — the gate is still real."

4 · 8-MINUTE TALK TRACK

Time	Where	Do / Say
0–1	Overview	RAG vs us card · Reset · Promise: catch → protect → prove
1–3	Knowledge	Ask policy Q · show file answer · no AI call
3–4	Human vault	Save fix · re-ask · vault hit
4–7	Outbound	Safe confirm · Medium rewrite · High private
7–8	Overview/Lab	Numbers · optional dry-run + evidence

5 · PILOT vs PRODUCTION

Topic	Pilot (today)	Production (later)
Gate decisions	Real Safe/Medium/High	Same + SIEM, SSO, policy packs
Outside AI	Groq live or simulate	Approved vendor under contract
Private AI	Laptop Ollama	GPU cluster / private cloud
Knowledge	Local files + vault JSON	Internal search + approved vault
Evidence	Files on this PC	Central audit + retention
Users	Demo room	Company-wide, managed

SAY: "Simulated reply ≠ fake security. Detection, hide, confirm, and High fail-closed are live."

6 · THREE PATHS

Uncertainty fails **up** (toward Medium/High) — never quietly down to Safe.

Path	When	Leaves?	Who answers?	Demo example
Safe	Clean text	Yes after confirm	Outside AI	Unit-test checklist
Medium	Business secrets (names, $)	Yes, cleaned	Outside AI	Deal $ / company name
High	Keys, IDs, PHI-class	No	Private AI only	API key / EMP + salary

Medium rewrite rule

- Same question meaning — not a new request.
- Real company → made-up company (never “our company”).
- Exact $ / % → softer wording (“tens of millions,” “sizable decline”).

7 · EVERY TAB

Tab	Purpose	Punch line
1 Overview	Scoreboard + charts	Catch → protect → prove
2 Knowledge	Policy Q from files/vault	No AI call on this tab
3 Outbound Gate	Free-form prompts	Safe / Medium / High live
4 Human vault	People fix gaps	Never trains a model
5 Pilot Lab	Dry-run + evidence	Prove without calling AI

8 · SCOREBOARD — READ THE NUMBERS

Number	Means
Could have leaked	Hide-path secrets in outbound prompts (synced to what we mask)
Sensitive in docs	Sensitive wording in local file answers (not sent to AI)
Latest leak risk %	Risk on last ask (local knowledge = no outside risk)
Outside AI skipped	Answered from files/human — no outside call
Human answers saved	Vault writes
Boundary split	Local vs external vs blocked
Lanes by team	Safe/Medium/High by team
Knowledge by team	Local / human-corrected / human gap
Sensitive by team	Last ask vs cumulative so far

9 · HIDE & REWRITE

Risky original	Safer same ask
Amazon Inc after a 30% drop?	Bluefinch Technologies after a sizable decline?
Flipkart Inc cutting costs 12%	Cedar Peak Labs cutting costs by a modest %
Sarah Chen · Contoso · $45M	Cameron Shaw · Ironwood · tens of millions

10 · PRIVATE AI — HONEST LINE

- High ≠ Safe. Safe = okay to send outside. High = kept inside.
- Ollama on this PC does not publish High prompts to the public (localhost).
- Private AI may still **see** the secret — containment, not “harmless.”
- Model saying “no” = safety filter. Win = refusal stayed on-box, not at a vendor.
- Private-only for everything is valid; hybrid is cost/quality after the gate.

SAY: “Private AI reduces who can steal the secret. It does not make pasting secrets into chat OK.”

11 · PROS · CONS · FIX IN PRODUCTION

Use this table when challenged. Own the con — then land the fix. That is what amazes rooms.

PRO — what works	CON — challenge	FIX IN PRODUCTION
Real gate before AI: detect → route → prove	Pattern detector can miss novel secret shapes	Add ML classifiers + continuous rule packs + human review queue for misses
High never goes to outside AI (fail-closed)	Private model may still see cleartext secrets	Always strip/replace secrets before private call; optional zero-LLM High playbooks
Medium keeps ask meaning with fictional names	Rewrite quality can feel stiff on edge cases	Enterprise rewrite model + human edit before send; A/B tune synonyms
Knowledge tab never invents policy via LLM	Gaps need humans — slower than auto-guess	SLA on vault queue; specialist owners; autocomplete from approved corpus
Pilot runs on one laptop — easy to demo	Not enterprise scale / SSO / HA yet	Gateway cluster, SSO, SIEM hooks, retention policies
Simulate mode = demo without vendor spend	Simulated prose can look “fake” if not disclosed	Always disclose mode; turn live Groq for executive rooms that need it
Local Ollama = no public egress on High	Small models refuse / weak answers on secrets	Private GPU fleet + enterprise system prompt or fixed rotate/revoke templates
Evidence locker for auditors	Local JSONL is not a compliance vault	Central immutable audit store, access control, legal hold

12 · KILLER Q&A — SAY THESE OUT LOUD

Q: How is this different from normal “docs + AI”?

A: Normal setups often send clear company text to a vendor AI. We add a control plane: detect secrets, hide before egress, High stays private, humans fill gaps — vault never trains a model. Show the Overview comparison card.

Q: Why use outside AI if we have private AI?

A: You can go private-only — most secure. Hybrid is optional: save private capacity for High; use stronger/cheaper outside AI only after the gate says clean or cleaned. ExhaustShield decides what may leave — it does not require leaving.

Q: Is High “safe”?

A: High is contained — not “Safe lane.” Safer against public leak. Best practice still: never paste live secrets. We catch, contain, and push safer asks.

Q: Will Ollama leak High to the public?

A: In this pilot, High calls localhost only. Nothing is sent to Groq for High. Production = private inference inside our network with the same rule.

Q: Why does private AI often say no?

A: Model safety training refuses key/password topics. That refusal happened on our machine — not at a public vendor. Production fix: playbook answers (rotate/revoke) without feeding raw secrets.

Q: What if detection misses something?

A: No system is 100%. We fail closed on uncertainty, re-scan before send, expand rules, keep evidence. We manage residual risk — we don’t pretend zero risk.

Q: Is simulate mode cheating?

A: No — we disclose it. Routing and redaction are live. Only outside prose may be stubbed without a key. Live Groq is one restart away.

Q: Will this slow people down?

A: Safe = one confirm. Medium = short clean-up. High = automatic private. Friction where risk is high; clean work stays fast.

Q: Who owns wrong answers?

A: Knowledge comes from files or humans. Gaps go to people. Vault is approved truth — not model invention.

Q: Build vs buy?

A: This pilot proves the pattern with working evidence. Rollout is IT + Security + Policy. You’re buying a control plane story you can demo today — not a slide.

Q: How does Groq fit?

A: Groq is the outside AI for Safe and cleaned Medium only — after confirm and last-mile check. High never goes to Groq.

Q: Does Reset wipe human work?

A: Scoreboard and evidence clear. Human vault and review queue stay — corrections survive demos.

13 · DEMO RECOVERY LINES

What goes wrong	Recovery line
High = model refusal	“Safety filter on-box. Important: it never left. Production = playbook answer.”
Ollama down	“Fail-closed. Still refused outside send. That’s a security win.”
Canned outside reply	“Simulate mode. Gate is real; prose is stubbed. We can go live.”
Stale numbers after reset	Ctrl+F5 · restart start_demo.bat if needed.
Weak rewrite	“Edit before send — human stays in control.”
They demand private-only	“Valid. Hybrid is optional. Gate supports either policy.”

14 · FILE GUIDE + WHO SPEAKS

When someone asks “where does this live?” — point to the file and the owner voice.

File / folder	Used for	Who speaks
static/index.html	All screens people see — tabs, charts, buttons	Product / Demo lead
gateway.py	Server: APIs, session stats, Reset Demo, wiring	Engineering
router.py	Safe/Medium/High decisions · Groq vs Ollama	Security + Eng
detector.py	Finds secrets · mask · rewrite hints	Security
knowledge.py	Tab 2 answers from company files	Knowledge / Ops
protection.py	Evidence / audit helpers for Pilot Lab	Security / Audit
start_demo.bat	One-click demo start · loads Groq settings	Demo lead
start_gateway.py	Starts the API server	Engineering
groq_key.local.bat	Live outside-AI key (local only — don’t share)	Demo lead / Eng
confidential_data/	Sample policy docs for Knowledge tab	Product / Knowledge
reviews/corrections.json	Human vault — approved answers	Ops / Knowledge
reviews/queue.json	Gaps waiting for a human	Ops
reviews/protection_audit.jsonl	Evidence locker records	Security / Audit
reviews/demo_session.json	Live scoreboard state for this run	Demo lead
build_ceo_playbook.py	Builds this PDF	You (presenter)
ExhaustShield_CEO_Demo_Playbook.pdf	This guide	You (presenter)

*TIP: Name-drop three files in the room — **router.py** (decide), **detector.py** (find secrets), **static/index.html** (prove on screen).*

15 · TODAY · NEXT

Ships in this pilot

- Three-lane Outbound Gate + confirm + last-mile re-scan
- Knowledge without LLM · Human vault · Pilot Lab evidence
- Overview analytics · rewrite with fictional names · expanded High rules
- Full Reset Demo · collapsible RAG vs us card

Credible next steps

- SSO · per-team policy · SIEM / tickets on High
- Private GPU inference · High playbooks without raw secrets
- Connectors for chat tools / coding assistants
- Admin rule tuning · retention · compliance export packs

16 · ONE-PAGE CHEAT SHEET

Start: start_demo.bat → http://127.0.0.1:8000/ → Reset Demo → Ctrl+F5 if needed.

Story: Catch → Protect (hide / private) → Prove (numbers + evidence).

Safe = confirm → outside · **Medium** = clean → outside · **High** = private only.

Honesty: High ≠ Safe · Private ≠ harmless · Simulate ≠ fake gate · Refusal ≠ public leak.

When challenged: Own the CON → land the FIX IN PRODUCTION (Section 11).

CLOSE: “People will use AI. ExhaustShield makes sure the risky stuff doesn’t leave quietly — and when we don’t know, a human answers, not a guessing model.”

— End — Re-read Sections 11 + 12 the night before any executive demo.