EXHAUSTSHIELD · BEGINNER HANDBOOK

Copilot Hybrid Setup

Users → Copilot Agent (Teams) → FastAPI → ExhaustShield engine

What this PDF is for
Follow every part in order to create the hybrid setup end to end: host the gateway, wire Copilot Studio, publish to Teams, and pin the 5-tab dashboard.

Golden rule
Never let Studio classify secrets. Always call `/api/generate` and trust the lane it returns.

Source markdown: **HANDBOOK_Copilot_Hybrid.md**
APIs aligned to this repo's **gateway.py**
Designed by Manikandan

Contents

1. What you are building
2. Licences checklist
3. Architecture and workflows
4. Lane formulas
5. Cost formulas
6. Part A — Prepare FastAPI on your server
7. Part B — Prove the three APIs
8. Part C — Create the Copilot Studio agent
9. Part D — Connect Studio to FastAPI
10. Part E — Build the topics
11. Part F — Publish to Teams
12. Part G — Pin the 5-tab dashboard
13. Part H — Test script
14. Part I — Rollout and go-live
15. Troubleshooting
16. Adaptive Card JSON
17. Questions for IT / admin

Users → Copilot Agent (Teams) → FastAPI → ExhaustShield engine

This handbook is written for a beginner. If you follow every step in order, you can stand up the hybrid setup end to end.

Designed for: ExhaustShield project in this folder

Pattern: Copilot Studio is only the front door. FastAPI (`gateway.py`) is the brain.

Golden rule: Never let Copilot Studio classify secrets. Always send the raw prompt to `/api/generate` and trust the lane it returns.

Table of contents

- **What you are building**
- **Licences checklist**
- **Architecture and workflows**
- **Lane formulas (how the engine decides)**
- **Cost formulas (simple math)**
- **Part A — Prepare the FastAPI engine on your server**
- **Part B — Prove the three APIs work**
- **Part C — Create the Copilot Studio agent**
- **Part D — Connect Studio to FastAPI**
- **Part E — Build the topics (menu + knowledge + gate)**
- **Part F — Publish to Teams**
- **Part G — Pin the full 5-tab dashboard**
- **Part H — Test script**
- **Part I — Rollout and go-live**
- **Troubleshooting**
- **Copy-paste Adaptive Card JSON**
- **Questions to ask IT / admin**

1. What you are building

End state

Piece	What the user sees	What actually runs
Copilot Agent in Teams	Chat menu: Policy / AI chat / Open dashboard	Studio topics + HTTP calls
FastAPI on company server	Invisible	**`gateway.py`** lanes + knowledge
Web dashboard	5 tabs in browser / Teams website tab	Same FastAPI + **`static/index.html`**

Flow (one picture)

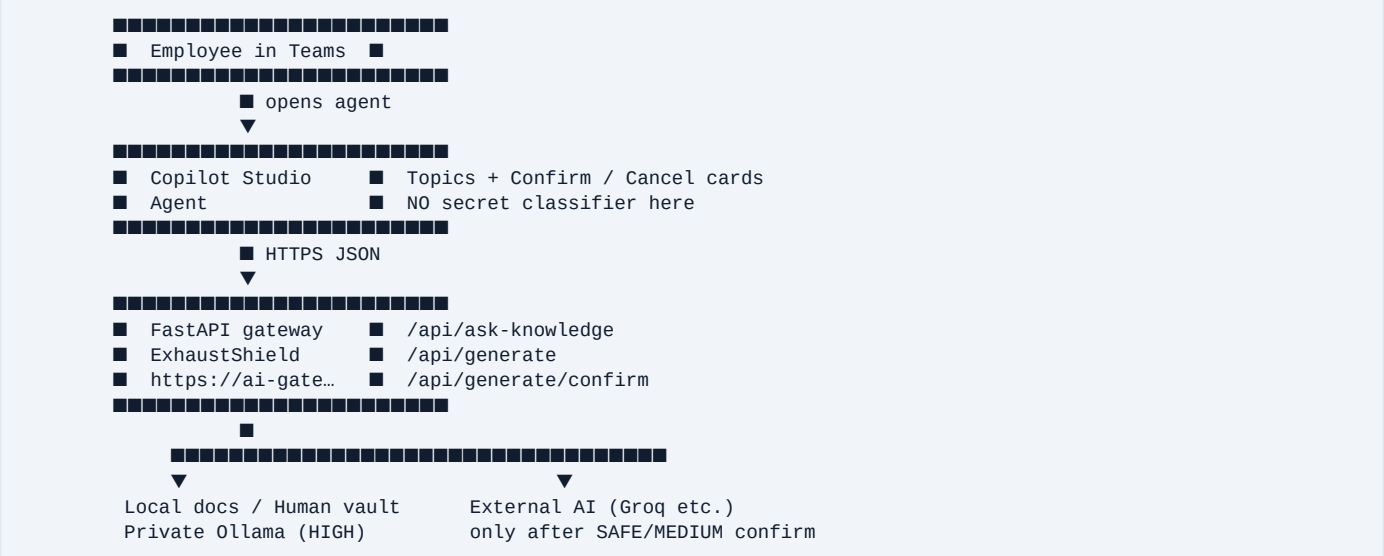

What this handbook does not do

- It does **not** rebuild Overview / Vault / Lab as Studio screens.
- It does **not** replace Microsoft Purview.
- It does **not** require buying M365 Copilot ($30) seats for everyone.

2. Licences checklist

Print this and tick with IT.

Required

#	Licence / access	Why	Do you need to buy new?
1	Microsoft 365 work account (E3 or E5 typical)	Teams + Entra sign-in	Usually already have
2	Access to a **Power Platform environment**	Where Studio agents live	Ask admin
3	**Copilot Studio** maker permission	Create and edit the agent	Ask admin (role / licence)
4	Permission to **publish to Microsoft Teams**	Users can open the agent	Admin approval often needed
5	Company server (or VM) you can run Python on	Hosts FastAPI	Prefer existing server
6	Outbound HTTPS from that server	Call Groq / Azure OpenAI	Network allowlist

Usage / consumption (not always a “seat”)

#	Item	When it costs money
7	**Copilot Studio credits / capacity**	Each conversation with *your* custom agent may consume credits (especially users without full M365 Copilot fair-use coverage). Ask admin for Dev trial + Prod PAYG or credit pack.
8	**Groq (or Azure OpenAI) tokens**	Only when SAFE/MEDIUM confirm sends text to external AI

Optional (not required for hybrid)

Item	Notes
M365 Copilot add-on $30 / user / month	For Word/Excel Copilot productivity — **not required** for this agent
Microsoft 365 E5 (vs E3)	Deeper Purview — complementary, not required for FastAPI gate
Purview Suite add-on (~$12 on E3)	Helps Copilot-side DLP — does not replace ExhaustShield lanes
Power Automate Premium	Only if you refuse HTTP/custom connector and use premium connectors

List prices to remember (USD, annual commit, Sep 2026 context)

SKU	Per user / month
Microsoft 365 E3	**$39** (from Jul 2026; was $36)
Microsoft 365 E5	**$60** (from Jul 2026; was $57)
M365 Copilot add-on	**$30**
ExhaustShield hybrid on existing server	**~$0 seats** + tokens + Studio credits

3. Architecture and workflows

Workflow A — Company knowledge (Tab 2 behaviour)

```
User types policy question
        ↓
Studio Topic "Knowledge"
        ↓
POST /api/ask-knowledge  { "question": "..." }
        ↓
FastAPI: vault → local near-match → human gap
        ↓
NEVER calls external LLM on this path
        ↓
Studio shows answer (+ sources if present)
```

Workflow B — Governed AI chat (Tab 3 behaviour)

```
User types free-form prompt
        ↓
Studio Topic "Gated chat"
        ↓
POST /api/generate  { "prompt": "..." }
        ↓
FastAPI detector + classify_route
        ↓
■■■
```

API `guardian.risk_level`	User-facing lane	What happens
`high`	HIGH	**Fail-closed** → private / blocked — **do not** call external AI from Studio

Confirm formulas (exact values your API accepts)

After **`/api/generate`** returns **`awaiting_confirm: true`** and a **`pending_id`**:

Situation	`POST /api/generate/confirm` body
SAFE — send as typed	`{ "pending_id": "<id>", "chosen": "as_is" }`
MEDIUM — placeholders	`{ "pending_id": "<id>", "chosen": "redacted" }`
MEDIUM — suggested rewrite	`{ "pending_id": "<id>", "chosen": "rewrite" }`
Cancel anything	`{ "pending_id": "<id>", "chosen": "cancel" }`

Alias: **`version`** is accepted instead of **`chosen`** (same values).

Knowledge formula

```
IF question empty → error
ELSE search vault / local policy docs
IF good local match → answer locally (sent_to_llm = false)
ELSE human gap / review path
NEVER silent fallback to Groq/Ollama on /api/ask-knowledge
```

5. Cost formulas (simple math)

Per chat (external model)

Assume one gated chat = **1,000 input + 500 output** tokens.

Groq gpt-oss-120b (list):

`$0.15 / 1M input` and **`$0.60 / 1M output`**

```
Cost_per_chat =
  (1000 / 1_000_000) * 0.15
+ (500  / 1_000_000) * 0.60
= 0.00015 + 0.00030
= $0.00045
```

Monthly token cost (existing server → hosting = $0)

```
Chats_per_month = Users × Chats_per_user_per_day × 22
Token_cost_month = Chats_per_month × 0.00045
Token_cost_year  = Token_cost_month × 12
```

Scenario	Chats / user / day	500 users / yr	3,000 users / yr
BEST	1	~$59	~$356
NORMAL	3	~$178	~$1,070
WORST	8	~$475	~$2,851

Add **Studio credits** on top (ask admin for your tenant's meter).

Do **not** add $30×users unless you separately buy full Copilot.

6. Part A — Prepare the FastAPI engine on your server

A1. Choose URL

Pick a hostname, for example:

```
https://ai-gate.yourcompany.com
```

You will use this everywhere in Studio.

A2. Copy project (no secrets)

Copy this folder to the server **except**:

- **`groq_key.local.bat`** (never copy keys into Git / tickets)
- Use **`groq_key.local.bat.example`** as template

A3. Install Python packages

On the server (PowerShell or bash):

```
cd "/path/to/Exhaust Shield"
python -m pip install -r requirements.txt
```

A4. Configure external AI

- Copy example → **`groq_key.local.bat`** (Windows) or set the same env vars in Linux service.
- Set:

```
EXTERNAL_AI_MODE=live          # or simulate for zero-cost pilot
EXTERNAL_AI_BASE_URL=https://api.groq.com/openai/v1
EXTERNAL_AI_API_KEY=gsk_...    # real key only on server
EXTERNAL_AI_MODEL=openai/gpt-oss-120b
```

- For a **$0 pilot**, use:

```
EXTERNAL_AI_MODE=simulate
```

A5. Start the gateway

Windows (from project folder):

```
python start_gateway.py
```

Default listen: **`http://127.0.0.1:8000`**

A6. Put HTTPS in front (IIS / nginx)

Reverse-proxy:

```
https://ai-gate.yourcompany.com  →  http://127.0.0.1:8000
```

Checklist:

- ■ TLS certificate valid
- ■ App runs as a service (survives reboot)
- ■ Firewall allows Power Platform / Teams egress to this host
- ■ Prefer auth (Entra / API key) before production

A7. Optional High lane (Ollama)

If HIGH should call a private model on the same box:

- Install Ollama
- Pull a model (e.g. **`llama3.2`**)
- Keep default **`OLLAMA_URL=http://127.0.0.1:11434/api/generate`**

If Ollama is missing, HIGH still **must not** go external — fail closed.

7. Part B — Prove the three APIs work

Use browser, Postman, or **`curl`**. Replace the base URL.

B1. Health

```
GET https://ai-gate.yourcompany.com/api/health
```

Expect JSON with **`"ok": true`**.

B2. Knowledge

```
POST https://ai-gate.yourcompany.com/api/ask-knowledge
Content-Type: application/json

{
  "question": "What is our leave policy?",
  "prefer_local_corrections": true
}
```

Expect fields like:

- **`ok`** (if present) / **`answer`** / **`sources`** / **`needs_human_review`**

B3. Generate (SAFE or pending)

```
POST https://ai-gate.yourcompany.com/api/generate
Content-Type: application/json

{
  "prompt": "Explain the difference between TCP and UDP in simple words"
}
```

Important response fields:

Field	Meaning
`awaiting_confirm`	**`true`** = nothing sent yet
`pending_id`	Pass this to confirm
`confirm_kind`	**`safe`** or **`moderate`**
`guardian.risk_level`	**`safe`** / **`moderate`** / **`high`**
`redacted_sent`	Placeholder version (MEDIUM)
`rewrite_suggestion`	Suggested rewrite (MEDIUM)
`response` / **final answer**	Often empty until confirm

B4. Confirm SAFE

```
POST https://ai-gate.yourcompany.com/api/generate/confirm
Content-Type: application/json

{
  "pending_id": "PASTE_ID_FROM_GENERATE",
  "chosen": "as_is"
}
```

B5. Confirm MEDIUM (hide secrets)

```
{
  "pending_id": "PASTE_ID",
  "chosen": "redacted"
}
```

B6. Cancel

```
{
  "pending_id": "PASTE_ID",
  "chosen": "cancel"
}
```

Do not continue to Studio until B1–B4 work.

8. Part C — Create the Copilot Studio agent

C1. Open Studio

- Go to https://copilotstudio.microsoft.com
- Sign in with work account
- Select **Dev** environment (ask admin for the name)

C2. Create agent

- **Create** → **New agent**
- Name: **`ExhaustShield Gate`**
- Description:

```
Governed company policy search and outbound AI chat.
All risk decisions are made by our ExhaustShield FastAPI gateway.
```

- Create the agent.

C3. Paste system instructions (important)

In agent instructions / overview instructions, paste:

```
You are the Teams front door for ExhaustShield.

Rules:
1. Never classify whether a prompt is secret, safe, or high risk yourself.
2. For policy questions, always call AskKnowledge (our API).
3. For free-form AI questions, always call Generate (our API) first.
4. If Generate returns awaiting_confirm=true, show Confirm/Cancel and wait.
5. Only call ConfirmGenerate after the user confirms.
6. If guardian.risk_level is high, tell the user nothing was sent externally.
7. Never invent scrubbed text. Only show redacted_sent or rewrite_suggestion from the API.
8. Offer a link to the full dashboard for Overview, Vault, and Pilot Lab.
```

C4. Turn off unsafe generative shortcuts

If the UI offers “generative orchestration” that answers without tools:

- Prefer **topics + actions** for Knowledge and Gate
- Or force “use tools / call connectors” for those intents

Beginners: build **explicit topics** (next sections). Do not rely on the model guessing.

9. Part D — Connect Studio to FastAPI

You have two beginner-friendly options. Start with **D-Option 1**.

D-Option 1 — HTTP request actions inside topics (fastest learning)

In each topic you will add **Call HTTP request** (name varies slightly by Studio version):

- Method POST
- URL full path
- Header **`Content-Type: application/json`**
- Body JSON

Skip to **Part E**.

D-Option 2 — Custom connector (better for enterprise)

- Open https://make.powerapps.com → same environment
- **Custom connectors** → **New** → **Blank**
- Connector name: **`ExhaustShieldAPI`**
- Host: **`ai-gate.yourcompany.com`**
- Base URL: **`/`**
- Scheme: HTTPS

Action 1 — AskKnowledge

- Operation ID: **`AskKnowledge`**
- Request: **`POST /api/ask-knowledge`**
- Body:

```
{
  "question": "string",
  "prefer_local_corrections": true
}
```

Action 2 — Generate

- Operation ID: **`Generate`**
- Request: **`POST /api/generate`**
- Body:

```
{
  "prompt": "string"
}
```

Action 3 — ConfirmGenerate

- Operation ID: **`ConfirmGenerate`**
- Request: **`POST /api/generate/confirm`**
- Body:

```
{
  "pending_id": "string",
  "chosen": "string"
}
```

- Auth: as IT requires (API key header recommended for v1)
- **Create connector** → **Test** each operation with real payloads from Part B
- In Copilot Studio → **Tools** / **Actions** → add **`ExhaustShieldAPI`** actions to the agent

10. Part E — Build the topics (menu + knowledge + gate)

Create **four** topics.

Topic E0 — Conversation Start / Main menu

Name: **`Main menu`**

Trigger: On conversation start (and a phrase trigger **`menu`**)

Message node:

```
Welcome to ExhaustShield

I can help with:
1) Company knowledge (policies — local answers)
2) Governed AI chat (scanned before anything leaves)
3) Open full dashboard (all 5 tabs)

What do you want to do?
```

Ask a question with options:

Option text	Navigate to
Company knowledge	Topic E1
Governed AI chat	Topic E2
Open full dashboard	Topic E3

Topic E1 — Company knowledge

Name: `Company knowledge`

Trigger phrases: `policy`, `company knowledge`, `handbook`, `leave policy`

Steps:

- **Ask:** `What is your policy question?`

Save answer → variable `UserQuestion` (text)

- **Call action / HTTP** `AskKnowledge`

Body:

```
{
  "question": "{UserQuestion}",
  "prefer_local_corrections": true
}
```

- **Parse response** into variables (names depend on Studio UI):

Variable	JSON path (typical)
`KnowledgeAnswer`	`answer` or `restored_response`
`NeedsHuman`	`needs_human_review`
`BestScore`	`best_score`

- **Condition:**
- If `NeedsHuman` is true →

`I could not find a solid local policy match. Please use Human vault in the full dashboard, or rephrase.`

- Else →

`Here is the local policy answer:`

then show `{KnowledgeAnswer}`

- **Message:**

`Tip: this path does not call an outside AI model.`

- **Go to** Main menu (or End)

Topic E2 — Governed AI chat (full gate workflow)

Name: `Governed AI chat`

Trigger phrases: `ask ai`, `chat`, `generate`, `outbound gate`

Step E2.1 — Collect prompt

Ask: `What do you want to ask the AI?`

Save → `UserPrompt`

Step E2.2 — Call Generate

HTTP / action `Generate`:

```
{
  "prompt": "{UserPrompt}"
}
```

Save into variables:

Variable	From response
`Awaiting`	`awaiting_confirm`
`PendingId`	`pending_id`
`ConfirmKind`	`confirm_kind`
`RiskLevel`	`guardian.risk_level`
`LaneBadge`	`guardian.lane_badge`
`Redacted`	`redacted_sent`
`Rewrite`	`rewrite_suggestion`
`RiskScore`	`risk_analysis.risk_score` (if present)
`DirectResponse`	`response` (may be empty)

Step E2.3 — Branch on risk

Condition group:

Branch HIGH (`RiskLevel` equals `high`)

Message:

```
HIGH lane — this prompt was held by ExhaustShield.
Nothing was sent to external AI.
Open the full dashboard Outbound Gate / Pilot Lab if you need the private-path details.
```

→ Main menu. **Do not call ConfirmGenerate.**

Branch SAFE (`RiskLevel` equals `safe` AND `Awaiting` is true)

Message:

```
SAFE lane — no secrets detected.
Risk score: {RiskScore}
Pending id: {PendingId}

Send this prompt to external AI as written?
```

Adaptive Card / choice:

- **Confirm send** → E2.4a
- **Cancel** → E2.4c

Branch MEDIUM (`RiskLevel` equals `moderate` AND `Awaiting` is true)

Message:

```
MEDIUM lane — sensitive spans found. Nothing sent yet.
Risk score: {RiskScore}

Hidden/placeholder version:
{Redacted}

Suggested rewrite:
{Rewrite}

How do you want to proceed?
```

Choices:

- **Send hidden version** → E2.4b with `chosen=redacted`
- **Send suggested rewrite** → E2.4b with `chosen=rewrite`
- **Cancel** → E2.4c

Branch other / error

Show raw error or **`Unknown or expired pending_id`** guidance → ask user to retry Generate.

Step E2.4a — Confirm SAFE

Call **`ConfirmGenerate`**:

```
{
  "pending_id": "{PendingId}",
  "chosen": "as_is"
}
```

Show final answer from confirm response (**`response`** / **`restored_response`** / message fields — inspect one real response and map the field once).

Step E2.4b — Confirm MEDIUM

```
{
  "pending_id": "{PendingId}",
  "chosen": "redacted"
}
```

or

```
{
  "pending_id": "{PendingId}",
  "chosen": "rewrite"
}
```

Show final answer.

Step E2.4c — Cancel

```
{
  "pending_id": "{PendingId}",
  "chosen": "cancel"
}
```

Message: **`Outbound cancelled — nothing sent to any AI.`**

Step E2.5 — Return to menu

Always offer **Main menu** after the flow.

Topic E3 — Open full dashboard

Name: **`Open full dashboard`**

Trigger: **`dashboard`**, **`overview`**, **`vault`**, **`pilot lab`**

Message:

```
Open the full ExhaustShield dashboard (all 5 tabs):

Overview · Company knowledge · Outbound Gate · Human vault · Pilot Lab

https://ai-gate.yourcompany.com/

Use this for live metrics, vault corrections, and Pilot Lab evidence.
```

(Replace with your real URL.)

11. Part F — Publish to Teams

F1. Publish the agent

- In Copilot Studio → **Publish**
- Wait until publish succeeds

F2. Turn on Teams channel

- **Channels → Microsoft Teams**
- Enable / connect
- Open Teams link / availability settings

F3. Availability stages

Stage	Who can use it
1	Only you
2	Security + 5 pilot users
3	One Team / channel
4	Org-wide (after tests)

F4. Admin approval

If users cannot see the app:

- Teams admin centre → Manage apps
- Find **`ExhaustShield Gate`**
- Allow for org / specific users

F5. First user open

Teams → Apps → search agent name → Add → chat **Main menu**

12. Part G — Pin the full 5-tab dashboard

This is how beginners get **all 5 tabs** without rebuilding them in Studio.

- Open the Team where the agent is used
- Channel → **+** → **Website**
- Tab name: **`ExhaustShield`**
- URL: **`https://ai-gate.yourcompany.com/`**
- Save

Now:

- Chat agent = Knowledge + Gate
- Website tab = Overview + Vault + Lab + full UI

13. Part H — Test script

Run these exact tests. Tick each box.

Knowledge

#	Input	Pass if
H1	“What is our leave policy?” (with a matching local doc uploaded)	Local answer, not external model
H2	Nonsense question with no policy	Human gap / no silent LLM invent

SAFE

#	Input	Pass if
H3	“Explain TCP vs UDP simply”	**`risk_level=safe`**, confirm required
H4	Confirm	External answer returns (or simulate answer)
H5	Cancel instead	Message says nothing sent

MEDIUM

#	Input	Pass if
H6	Prompt including a fake deal value / employee salary style text your detector flags	**`moderate`**, shows redacted + rewrite
H7	Confirm **`redacted`**	Answer returns; outbound was scrubbed
H8	Cancel	Nothing sent

HIGH

#	Input	Pass if
H9	Prompt with API key / password / obvious secret	**`high`**; Studio never calls confirm-to-external
H10	Check server audit / Overview traffic	Shows High / private path

Hybrid plumbing

#	Check	Pass if
H11	User with only basic Copilot Chat (no $30 seat)	Can open your agent (if Studio licensing allows)
H12	Website tab	All 5 tabs load
H13	**`/api/health`**	**`ok: true`**

14. Part I — Rollout and go-live

Phase 0 — Builder (you)

- ■ FastAPI live on HTTPS
- ■ APIs tested with curl/Postman
- ■ Agent topics work for you

Phase 1 — Pilot (20–50 people, 30–60 days)

- ■ **`EXTERNAL_AI_MODE=simulate`** or free tier if possible
- ■ Budget cap suggestion: **$500** total
- ■ Collect: High holds, Medium scrubs, user complaints

Phase 2 — Production

- ■ **`EXTERNAL_AI_MODE=live`** + real key in server secret store
- ■ Studio Prod environment + credit plan agreed
- ■ InfoSec sign-off on HIGH never in external logs

Success metrics (90 days)

```
Adoption %     = Users who opened agent or dashboard / invited users
Gate coverage  = Gated chats / (gated chats + known shadow AI incidents)
High holds     = Count of HIGH lanes
Medium scrubs  = Count of moderate confirms with chosen=redacted|rewrite
Cost           = Groq invoice + Studio credits
```

15. Troubleshooting

Problem	Likely cause	Fix
Studio HTTP 401/403	Auth missing	Add API key / Entra to connector

Problem	Likely cause	Fix
Studio HTTP timeout	Network blocked	Allow Power Platform → ai-gate
`Unknown or expired pending_id`	Confirm too late or server restarted	Pending store is in-memory — re-run Generate; for production use sticky single instance
SAFE never asks confirm	You showed **`response`** early	Always check **`awaiting_confirm`**
HIGH still "answers" from Studio model	Generative orchestration answered alone	Disable free answer; force tool call
Knowledge invents policy	Wrong endpoint or generative fallback	Must call **`/api/ask-knowledge`** only
Teams users can't find agent	App not approved	Teams admin allow app
Cost spike	Studio credits + heavy chat	Cap pilots; check credit report

Production note: **`pending_id`** values live in process memory (**`pending_moderate`**). Run **one** gateway instance (or later add shared cache) so Confirm always hits the same process that created the pending id.

16. Copy-paste Adaptive Card JSON

Use in Topic E2 for SAFE confirm (paste into Adaptive Card node if available).

```
{
  "type": "AdaptiveCard",
  "version": "1.5",
  "body": [
    {
      "type": "TextBlock",
      "text": "ExhaustShield — SAFE confirm",
      "weight": "Bolder",
      "size": "Medium"
    },
    {
      "type": "TextBlock",
      "text": "No secrets detected. Send this prompt to external AI as written?",
      "wrap": true
    },
    {
      "type": "TextBlock",
      "text": "Pending: ${PendingId}",
      "wrap": true,
      "size": "Small"
    }
  ],
  "actions": [
    {
      "type": "Action.Submit",
      "title": "Confirm send",
      "data": { "choice": "as_is" }
    },
    {
      "type": "Action.Submit",
      "title": "Cancel",
      "data": { "choice": "cancel" }
    }
  ]
}
```

MEDIUM card — same pattern with three actions: **`redacted`**, **`rewrite`**, **`cancel`**.

17. Questions to ask IT / admin

Copy-paste into an email:

```
Subject: ExhaustShield hybrid — access needed

1) Which Power Platform environment should I use for Dev and Prod?
2) Do I have Copilot Studio maker rights? If not, which licence/role?
3) How are Copilot Studio credits billed for a custom agent in Teams?
4) Can Power Platform reach https://ai-gate.yourcompany.com?
5) Preferred API auth for the connector (API key vs Entra)?
6) Can we publish a custom Copilot Studio agent to Teams for a pilot group?
7) Can we add a Website tab in Team <NAME> to https://ai-gate.yourcompany.com/?
8) Who owns TLS + Windows/Linux service for the FastAPI process?
9) Is Groq (or Azure OpenAI) approved for SAFE/MEDIUM outbound after scrubbing?
10) Do we need InfoSec test that HIGH prompts never appear in external provider logs?
```

Quick end-to-end checklist (one page)

- ■ Server runs ExhaustShield on HTTPS
- ■ **`/api/health`**, **`/api/ask-knowledge`**, **`/api/generate`**, **`/api/generate/confirm`** work
- ■ Copilot Studio agent created with “never classify yourself” instructions
- ■ HTTP or custom connector points at the three APIs
- ■ Topics: Main menu, Knowledge, Gated chat (SAFE/MEDIUM/HIGH branches), Dashboard link
- ■ Publish → Teams channel enabled → admin approved
- ■ Website tab pinned for full 5-tab UI
- ■ Test script H1–H13 passed
- ■ Pilot users invited; credits + token budget agreed
- ■ Production go-live only after HIGH fail-closed proof

Remember for the board (one paragraph)

We keep Microsoft Copilot Chat and Purview for Office AI. ExhaustShield hybrid puts a **Copilot Agent in front of our FastAPI gate** so staff work in Teams, while **our engine** still decides SAFE / MEDIUM / HIGH / LOCAL. Full five tabs stay on the web dashboard. On an existing server, new cash cost is mainly **model tokens** (about **$0.00045 per gated chat**) plus **Studio credits** — not $30 Copilot seats for everyone.

Handbook aligned to this repo’s **`gateway.py`** *endpoints and confirm contract (*`pending_id` + `chosen`*). Designed by Manikandan / ExhaustShield pilot materials.*